

# Costumer\_churn

Loading libraries.

```
library(caret)
library(C50)
library(dplyr)
data(churn)
```

```
prop.table(table(churnTrain$churn))
```

```
##
##      yes      no
## 0.1449145 0.8550855
```

Notice that just 14% of customers will churn.

## Create train test indices for cross-validation

```
set.seed(42)
myFolds <- createFolds(churnTrain$churn, k = 5)
```

## Compare class distribution

```
i <- myFolds$Fold1
table(churnTrain$churn[i])
```

```
##
## yes  no
##  96 570
```

## Create myControl

. to have exact same cross validation folds for each model. This allow us to compare these model and make a fair comparison.

```
myControl <- trainControl(summaryFunction = twoClassSummary,
                           classProbs = T,
                           verboseIter = F,
                           savePredictions = T,
                           index = myFolds)
```

## Fit the models

1. glmnet model

```
set.seed(42)
model_glmnet <- train(churn ~ .,
                      data = churnTrain,
                      metric = "ROC",
                      method = "glmnet",
                      tuneGrid = expand.grid(
                        alpha = 0:1,
                        lambda = seq(.0001, 1, length = 20))
```

```

    ),
    trControl = myControl
  )
# Find the best ROC value
model_glmnet$results %>%
  filter(ROC == max(ROC))

##   alpha   lambda      ROC      Sens      Spec      ROCSD      SensSD
## 1      0 0.2106053 0.7686188 0.01397223 0.9996491 0.01167738 0.003908229
##           SpecSD
## 1 0.0004804584

#glmnet_pred <- predict(model_glmnet, newdata = churnTest)
#confusionMatrix(glmnet_pred, churnTest$churn)
#library(Metrics)
#auc(actual = ifelse(churnTest$churn == "yes", 1, 0), predicted = glmnet_pred[, "yes"])

```

Notice that  $\alpha = 0$  and  $\lambda = 0.2106053$  give the best result “ROC” = .7686188.

## 2. glm model

```

set.seed(42)
model_glm <- train(churn ~ .,
  data = churnTrain,
  method = "glm",
  metric = "ROC",
  trControl = myControl)

model_glm$results

##   parameter      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 1      none 0.7077977 0.2981229 0.9401754 0.03284925 0.02928006 0.01214128

#glm_pred <- predict(object = model_glmnet, newdata = churnTest)
#confusionMatrix(glm_pred, churnTest$churn)

#library(caTools)
#colAUC(glm_pred, churnTest$churn, plotROC = T)

```

3. Random Forest RF is slower to fit than glmnet but often(not always) more accurate than glmnet, easier to tune, captures threshold effect and variable interactions systematically.

```

set.seed(42)
model_rf <- train(churn ~ .,
  data = churnTrain,
  method = "ranger",
  metric = "ROC",
  tuneGrid = expand.grid(mtry = seq(4, ncol(churnTrain) * 0.8, 2),
    splitrule = "gini"),
  trControl = myControl)

#library(dplyr)
model_rf$results %>%
  filter(ROC == max(ROC))

##   mtry splitrule      ROC      Sens      Spec      ROCSD      SensSD
## 1    16      gini 0.903621 0.5418926 0.9891228 0.007002841 0.0619377

```

```
##          SpecSD
## 1 0.003013257

#rf_pred <- predict(object = model_rf, newdata = churnTest)
#confusionMatrix(rf_pred, churnTest$churn)
```

#### 4. Gradient Boosting

```
set.seed(42)
model_xgb <- train(churn ~ .,
  data = churnTrain,
  method = "xgbTree",
  metric = "ROC",
  tuneGrid = expand.grid(eta = .01,
    gamma = 0,
    max_depth = c(5, 10),
    colsample_bytree = 1,
    min_child_weight = 1,
    subsample = .75,
    nrounds = seq(100, 200, 50)),
  trControl = myControl
)

model_xgb$results %>%
  filter(ROC == max(ROC))

##      eta max_depth gamma colsample_bytree min_child_weight subsample nrounds
## 1 0.01         5      0              1              1      0.75      200
##      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 1 0.9064618 0.6308993 0.9834211 0.00666029 0.06459166 0.004095655
```

## Comparing models

Make sure they were fit on the same data

Selection criteria .Highest average AUC .Lowest standard deviation AUC A. Make a list of models

```
model_list <- list(glmnet = model_glmnet,
  rf = model_rf,
  xgb = model_xgb,
  glm = model_glm
)
```

*#Collect resamples from the cv folds*

```
set.seed(42)
resamps <- resamples(model_list)
resamps
```

```
##
## Call:
## resamples.default(x = model_list)
##
## Models: glmnet, rf, xgb, glm
## Number of resamples: 5
## Performance metrics: ROC, Sens, Spec
```

```
## Time estimates for: everything, final model fit
```

```
#Summarize the results
```

```
summary(resamps)
```

```
##
```

```
## Call:
```

```
## summary.resamples(object = resamps)
```

```
##
```

```
## Models: glmnet, rf, xgb, glm
```

```
## Number of resamples: 5
```

```
##
```

```
## ROC
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
## glmnet	0.7530372	0.7621334	0.7718537	0.7686188	0.7720012	0.7840685	0
## rf	0.8956544	0.9000341	0.9013845	0.9036210	0.9073745	0.9136577	0
## xgb	0.8953476	0.9056080	0.9085941	0.9064618	0.9110159	0.9117435	0
## glm	0.6528747	0.7094009	0.7121057	0.7077977	0.7260767	0.7385306	0

```
##
```

```
## Sens
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
## glmnet	0.01036269	0.0129199	0.01295337	0.01397223	0.01295337	0.02067183
## rf	0.46113990	0.5077720	0.53488372	0.54189260	0.59326425	0.61240310
## xgb	0.54145078	0.6010363	0.64766839	0.63089931	0.64857881	0.71576227
## glm	0.27461140	0.2772021	0.28165375	0.29812293	0.31606218	0.34108527

```
## NA's
```

```
## glmnet 0
```

```
## rf 0
```

```
## xgb 0
```

```
## glm 0
```

```
##
```

```
## Spec
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
## glmnet	0.9991228	0.9991228	1.0000000	0.9996491	1.0000000	1.0000000	0
## rf	0.9846491	0.9881579	0.9890351	0.9891228	0.9916667	0.9921053	0
## xgb	0.9771930	0.9815789	0.9846491	0.9834211	0.9868421	0.9868421	0
## glm	0.9263158	0.9311404	0.9421053	0.9401754	0.9438596	0.9574561	0

```
bwplot(resamps, metric = "ROC")
```

